



CPEN 311: Digital Systems Design

**On-Chip Memories**  
Memory Hierarchies and Scheduling

## Learning Objectives for Slide Set

---

After this slide set, you will be able to:

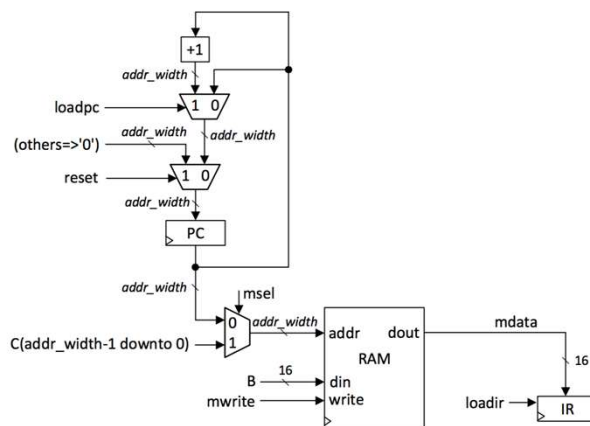
1. Understand the differences, advantages and disadvantages of off-chip memory and on-chip memory
2. Understand three ways on-chip memory can be implemented in an FPGA
3. Understand the notion of memory ports, and be able to speak about single-port, dual-port and true-dual port memory configurations
4. Understand how the number and type of ports impacts overhead
5. Be able to schedule an algorithm subject to the availability of memory ports
6. Understand typical on-chip memory timing
7. Be able to write Verilog descriptions containing memory

Page 2

## Memories

In CPEN211, you learned about memories, including address ports, data ports, write enable.

Depending on when you took CPEN 211, you might have implemented a processor containing a memory, like this:



Our focus will be different: using memories to implement algorithms

- They dictate scheduling of operations (1 access per cycle per port)
- Often, memory are the limiting factor in performance

Page 3


## TL;DR

be very careful about modules that use memory

- can't write to same port multiple times
  - if you do and use implicit memory instances, you will get "memories" from FFs
- can't load (these) memories within one cycle
- need to design your memory client carefully

Page 4

## On-Chip vs. Off-Chip



Local storage on-chip

For larger datasets, use on-board SDRAM

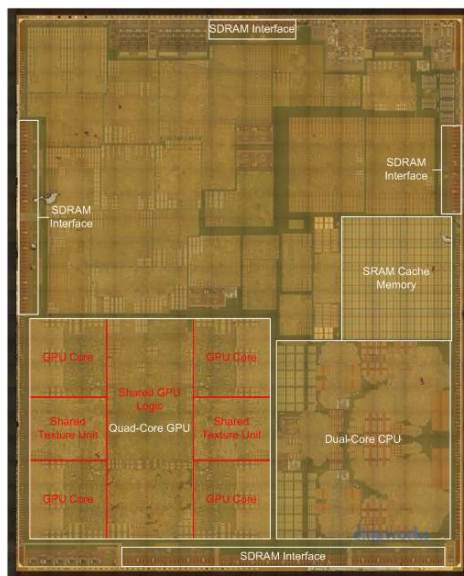
For much larger storage, connect external memories, or interface to internet (and use cloud storage)

In this slide set, we will focus mostly **on on-chip** memory


Page 5

## On-Chip (Embedded) Memory

In a custom chip, you can instantiate a memory block that is exactly the size you need. Often done using a **memory generator**



A8 Processor




A9 Processor

Page 6

## On-Chip (Embedded) Memory

In an FPGA, you need to use pre-fabricated memory resources



Page 7

Next topic: Ways to implement storage on an FPGA


- A) Embedded Memory Blocks
- B) Flip-Flops in the logic fabric
- C) Configuration Storage (only on certain FPGA devices, not ours)

This is important because these methods have very different overheads and constraints.



A) First way is Embedded Memory Blocks:

- FPGA Manufacturer includes memory blocks (at the expense of logic blocks):



Page 9

## Embedded Memory Blocks:

Altera Cyclone V SE (DE1-SoC board) and V E (DE0-CV board)

Each embedded memory block has 10Kbits

Device	M10K Blocks	Total Memory Bits
5CSEMA2	140	1,433,600
5CSEMA4	270	2,764,800
5CSEMA5	397	4,065,280
5CSEMA6	557	5,703,680

Device	M10K Blocks	Total Memory Bits
5CEBA2	176	1,802,240
5CEBA4	308	3,153,920
5CEBA5	446	4,567,040
5CEBA7	686	7,024,640
5CEBA9	1220	12,492,800

### **Advantages of using embedded memory blocks:**

- More efficient for large memories
  - In our chip up to 484kB / 4Mb on-chip storage
- Address decoders, muxes, etc. already implemented as part of the memory blocks (don't need to implement these using logic blocks)

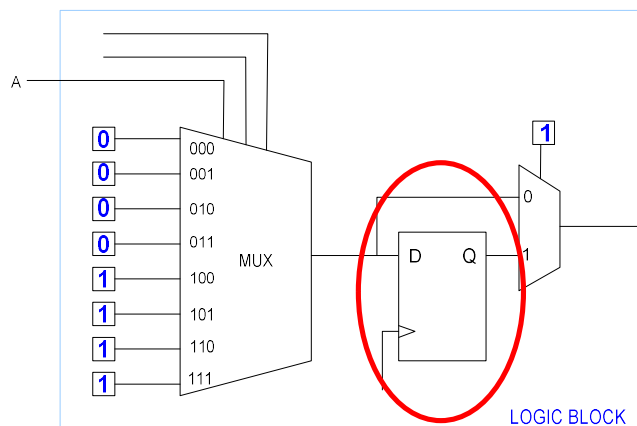
### **Disadvantages of using embedded memory blocks:**

- Less flexible
- Registers are faster for small sizes

For > 1kb, embedded memories are the way to go

## Flip-Flop Memory:

B) Second way to implement memory in an FPGA:  
Each logic element has a flip-flop



DE1-SoC (5CSEMA5) has 128k flip-flops

- Not a very efficient way to implement large memories!

**Advantages of using flip-flops to implement memory:**

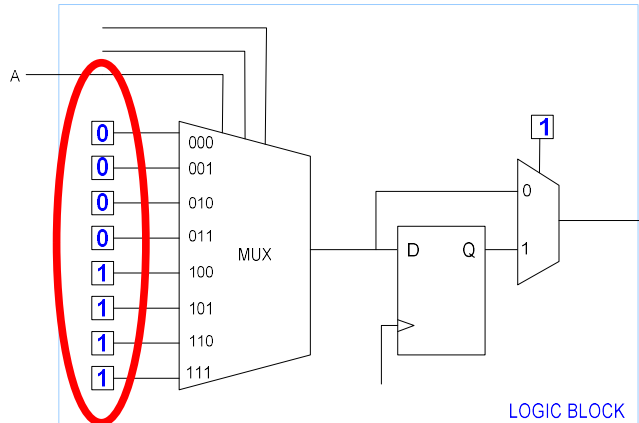
- Very flexible, you can combine registers however you like
- Fast connections between logic and memory, since registers are in the logic fabric

**Disadvantages of using flip-flops to implement memory:**

- Each bit of memory takes a logic block (quickly run out)
- You need to construct address decoders, output muxes, etc., out of general purpose logic. These also are large and will quickly fill up your chip (more on this coming...)

For <1kb, registers are the way to go.

C) Third way to implement memory in an FPGA is configuration storage:



Configuration bits are actually memory bits! Some FPGAs (not ours) lets you use these as user memory.

We won't talk about this option any more in this slide set.

Page 14

Next topic: Read and Write Port Configurations

- a) Fixed Read and Write ports in Embedded Memories
- b) Creating Read and Ports using Flip-Flop Memories

This is important because the read and write port configuration impacts how you can schedule your algorithm in hardware.

Page 15

## Embedded Memory Blocks

### Read and Write Ports


Page 16



## Cyclone II / V Memory Blocks

Memory blocks are highly configurable

Physically they are built like this:



Ports A & B share one physical memory, mild constraints on crossing  
Logically, they can be configured into several modes (see next slide)

Page 17

## Cyclone II / V Memory Blocks

An embedded memory block can be configured as one of:

Single-port memory	One Read or Write per clock cycle
Simple Dual-port memory	Simultaneous Read and Write.
True dual-port memory	Two operations, each can be read or write
ROM	A MIF initializes the ROM contents of these blocks.
FIFO buffers	FIFO Queue.

Cyclone II: M4K ( $512 \times 9 = 4,608$  bits = 4.5k bits)

Cyclone V: M10K ( $1024 \times 10 = 10,240$  bits = 10k bits)


Other FPGAs have similar memory structures

Important point: these memory blocks are very flexible.

- Use the correct mode for your application to make the most efficient use of blocks.

Page 18

## Single-Port Memory: Generic View

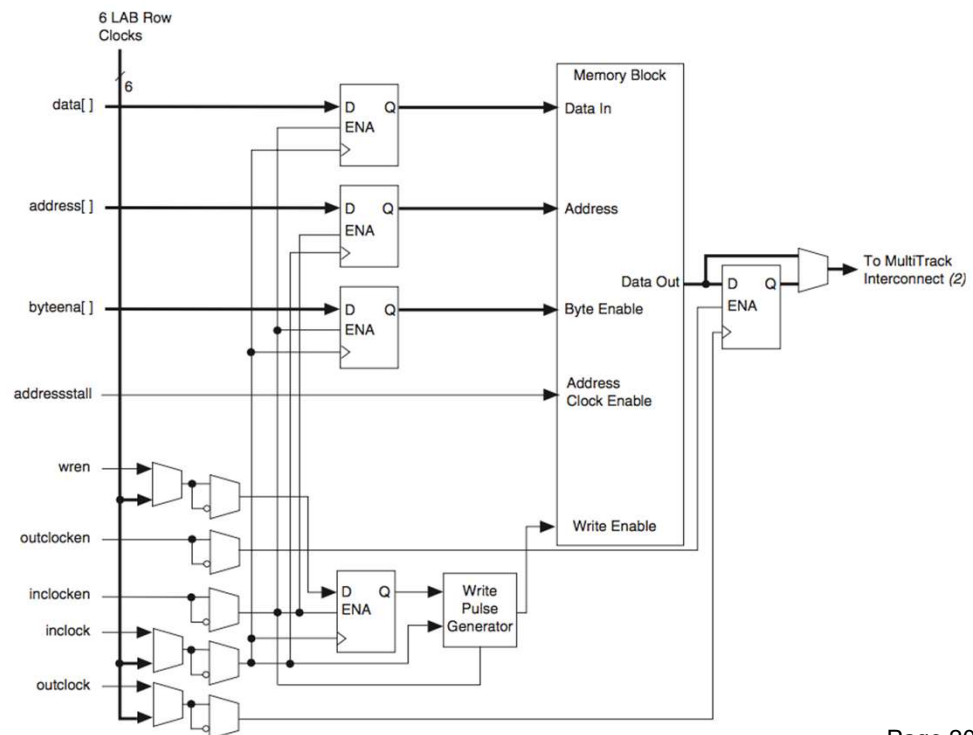


Each cycle, you can either write or read a single word from memory (wren determines whether you are writing or reading).

Write: put address on address bus, data on data bus and assert wren

Read: put address on address bus, de-assert wren, data appears on q

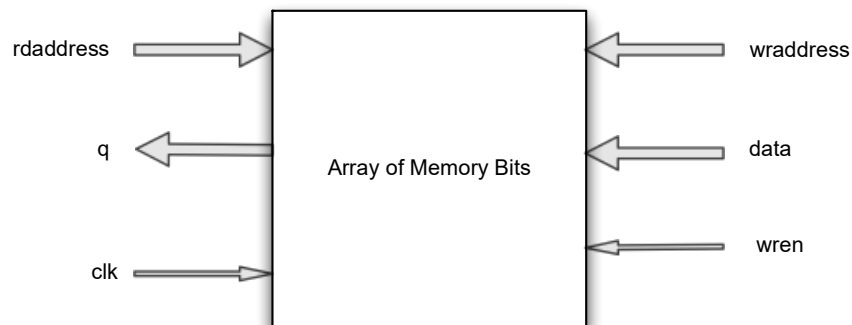
## Single-Port Memory: details



From Altera, Cyclone II handbook

Page 20

## Simple Dual-Port Memory: Generic View



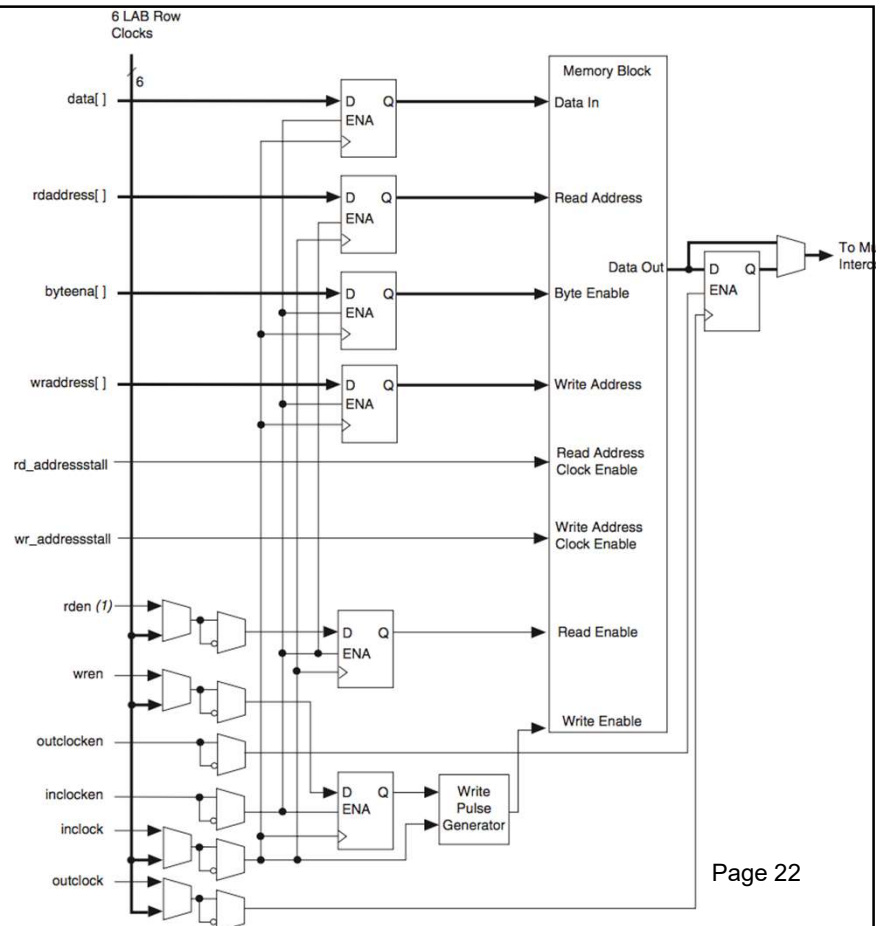
Each cycle, you can write a word *and* read a word

- Normally, addresses will be different
- If you write and read the same address at the same time, there are several possible policies (read old, read new)
- Some memories allow for separate clocks

This sort of memory is good for implementing FIFOs.


Page 21

## Simple Dual-Port Memory: details



From Altera, Cyclone II handbook

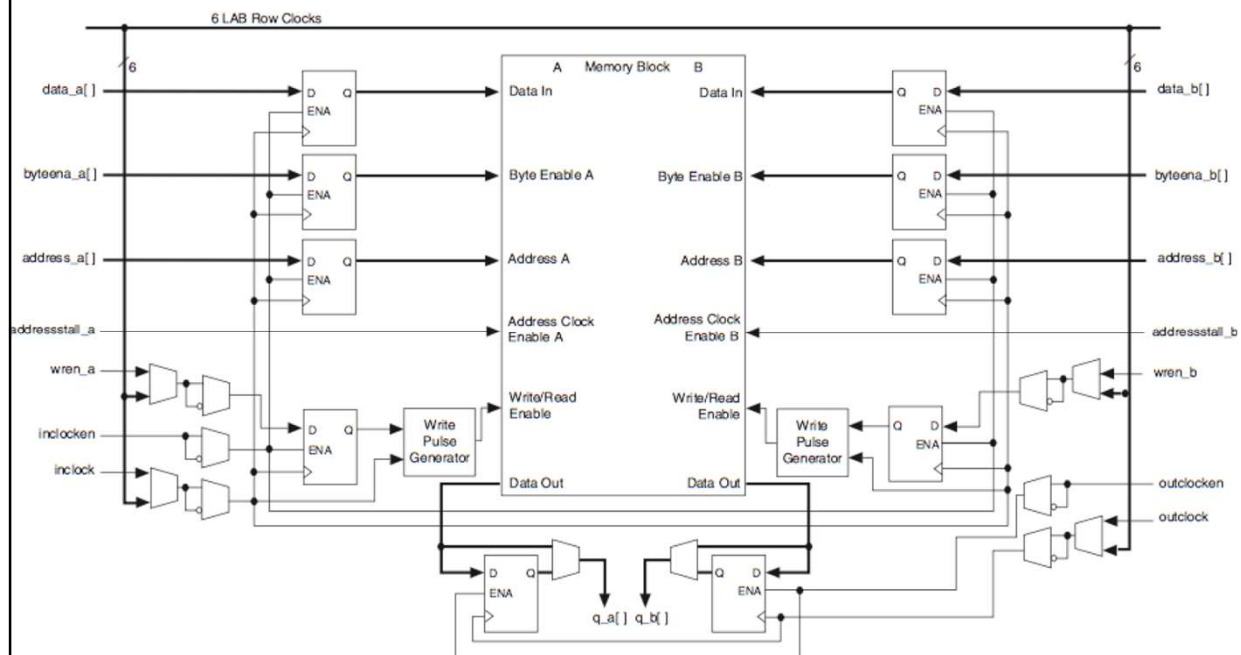
## True Dual-Port Memory: Generic View



Each cycle, you can perform two access, either of which can be a write or read

- Some systems allow two different clocks

## True Dual-Port Memory: details



From Altera, Cyclone II handbook

Page 24



When you use embedded memory blocks, this port circuitry is included in each memory block already.

What if we want to use flip-flop memory rather than the embedded blocks?

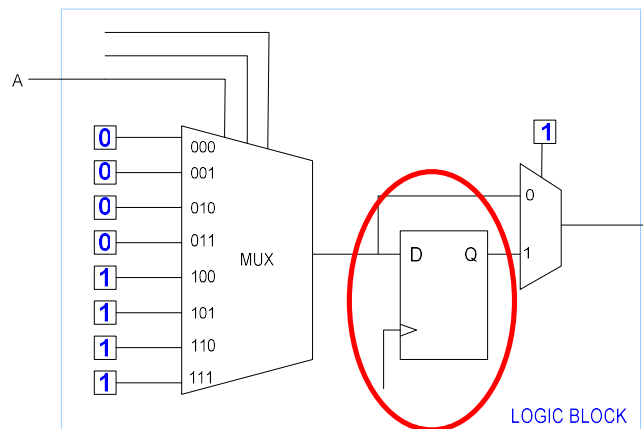
## Flip-Flop Memory

### Read and Write Ports

Page 26

## Flip-Flop Memory: Ports

What if you want to use flip-flop memory?




Read and write ports need to be implemented out of general purpose logic

**GOOD News:** you can implement exactly the port structure you want.

**BAD News:** ports use a lot of logic, very inefficient use of your chip

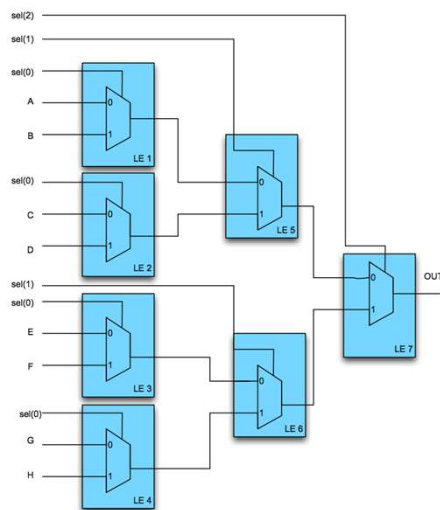
## Flip-Flop Memory: Read Ports

Consider implementing a read-port out of general-purpose logic:



How big is this mux?

Consider building an 8-1 mux



In this case, you could do 8 input, 1 bit mux in 7 logic blocks

$$4+2+1 = 7$$

In general, you can do a n input, 1 bit mux in n-1 logic blocks

$$(n/2) + (n/4) + \dots 1 = n-1$$


If the mux is m bits wide, you would need about  $m \cdot (n-1)$  logic blocks

Example: a 256 bit x 8 bit memory would need 2040 logic blocks just to implement one read port!

## Flip-Flop Memory: Write Ports

The size of each write port is a bit more tricky to work out, but we can approximate it:

- The heart of the write port is a decoder with  $n$  outputs. Each output goes high when the corresponding word is being written to



This decoder would need at least  $n$  logic blocks, since each logic block has only one output.

So a 256x8 bit memory would need at least 256 logic blocks to implement the write port (actually more than this)

## Flip-Flop memories use a lot of logic blocks:

Consider a 256 x 8 bit memory with one read port and one write port:

Logic blocks required to implement storage:  $256 \times 8 = 2048$

Logic blocks required to implement read port: 2040

Logic blocks required to implement write port: 256

Total is 4344 logic blocks. FPGA on DE2 has only has 33,216 logic blocks!

We can add more ports. Suppose we want 4 read ports (because we want to read 4 words each cycle):

$$2048 + 4 \times 2040 + 256 = 10464 \text{ logic blocks}$$

And this is a small memory!

Page 31

**Moral:** using flip-flop memories uses a lot of logic blocks.

**Danger:** Suppose you decide you want to use embedded memories.  
It turns out that if you don't write your Verilog "just right", the  
synthesis tool will give you flip-flop memories.

Why is this? That is the next topic.

Page 32



Next topic: Describing Memories in Verilog

- a) Explicit Instantiation
- b) Inferred Memories

Page 33

## Including Memories in Verilog

---

### 1. Explicit Instantiation:

- Explicitly indicate that you are including a memory using a pre-defined Altera macro. You can specify the number of ports, size, width, output registers, etc.
- Can be done either directly with a port map or through a Wizard

### 2. Inferred Memories:

- Write the behaviour of a memory using behavioural code and expect the tool to figure out that you want a memory (use the “array” construct in Verilog)
- Inferring to embedded RAMs is tricky: you have to get it right
- Inferring to flip-flop memories is relatively easy

## Explicit Representation

---

Most systems contain macros that you can use to explicitly indicate a memory.

In Altera's case, you can use ALTSYNCRAM

Details on ALTSNYCRAM are at

[http://quartushelp.altera.com/14.1/mergedProjects/hdl/mega/mega\\_file\\_altsynch\\_ram.htm](http://quartushelp.altera.com/14.1/mergedProjects/hdl/mega/mega_file_altsynch_ram.htm)

You can also use LPM\_RAM\_DQ, however, it has been depreciated (still available for backwards compatibility).

## ALTSYNCRAM (Verilog)

```
altsyncram altsyncram_component (  
    .address_a (address),  
    .clock0 (clock),  
    .data_a (data),  
    .wren_a (wren),  
    .q_a (sub_wire0));  
  
defparam  
    altsyncram_component.clock_enable_input_a = "BYPASS",  
    altsyncram_component.clock_enable_output_a = "BYPASS",  
    altsyncram_component.intended_device_family = "Cyclone V",  
    altsyncram_component.lpm_hint = "ENABLE_RUNTIME_MOD=YES,INSTANCE_NAME=S",  
    altsyncram_component.lpm_type = "altsyncram",  
    altsyncram_component.numwords_a = 256,  
    altsyncram_component.operation_mode = "SINGLE_PORT",  
    altsyncram_component.outdata_aclr_a = "NONE",  
    altsyncram_component.outdata_reg_a = "UNREGISTERED",  
    altsyncram_component.power_up_uninitialized = "FALSE",  
    altsyncram_component.ram_block_type = "M10K",  
    altsyncram_component.read_during_write_mode_port_a =  
"NEW_DATA_NO_NBE_READ",  
    altsyncram_component.widthad_a = 8,  
    altsyncram_component.width_a = 8,  
    altsyncram_component.width_byteena_a = 1;
```

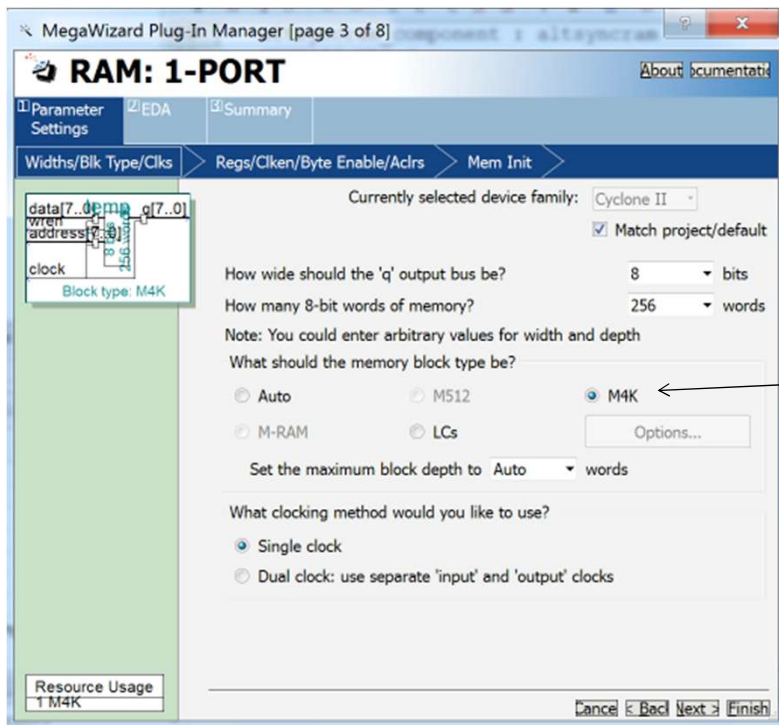
Signals to connect to the memory

Parameters that describe the type of memory you want

Page 36

## How to generate these?





Choose whether you want an embedded memory or flip-flop memory

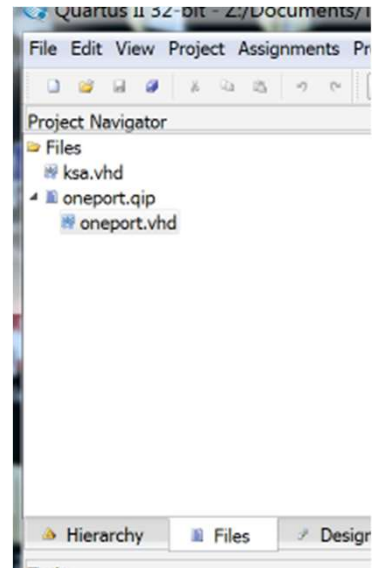
Goes through 8 pages.... at one point you can indicate VHDL or Verilog. Doesn't matter which

Page 38


Creates a lower level file with a module that you can include in your design:

```
module oneport (
    address,
    clock,
    data,
    wren,
    q);

input [7:0] address;
input clock;
input [7:0] data;
input wren;
output [7:0] q;
```



## Using this component



In your datapath or state machine, you can do a **write** as follows:


```
address <= 8'b011011000;  
data <= 8'b100111000;  
wren <= 1'b1;
```

Timing is very important

Page 40



## Using this component



In your datapath or state machine, you can do a **read** as follows:

```
address <= 8'b011011000;  
wren <= 1'b0;  
wait for up to two cycles  
value <= q;
```

Timing is very important

Page 41

In general, the more ports, the more simultaneous reads and writes you can do.

Should you always just use dual-port memories?

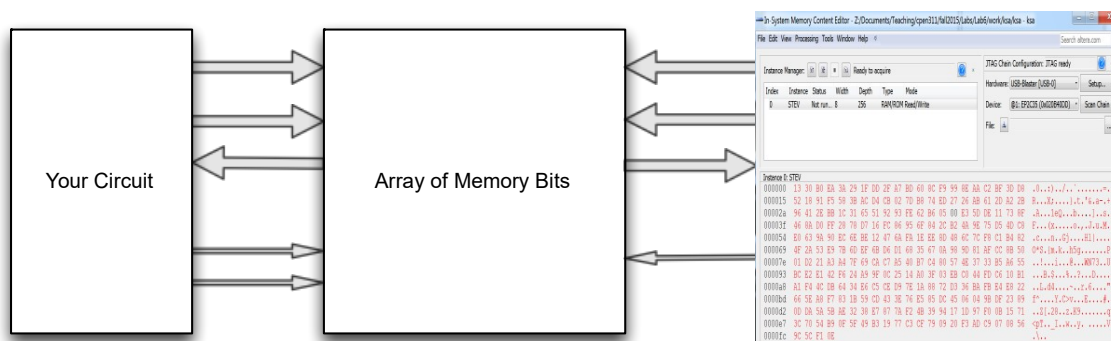
Not necessarily:

1. In a custom chip, dual-port memories are a bit slower and larger (not an issue in an FPGA because they are already “there”)
2. You might want to use the other port for something else
3. In some algorithms, it just might not help...

Page 42

Using other port for something else:

In the lab, you might find it useful to connect one port to the “In-System Memory Content Editor”...



Useful for debugging because you can see what is in your memory  
(also when we test your implementation!)

More details on this tool in the lab text

You can have as many memories in your system as you like  
Accesses to memories can be in parallel

Often, each major data structure is mapped to a different memory

Page 44

Important take away: limited read and write ports adds constraints for scheduling. Must take these into account when designing a circuit that talks to memory.

Page 45

## Inferred Memories

---

In Verilog, you can use the array construct to describe a memory

```
reg [7:0] mem [127:0];
```

You can then describe a read operation behaviourally:

```
some_variable <= mem[i];
```

You can describe a write operation behaviourally:

```
mem[i] <= some value
```

where i is a variable or signal and contains the address you want to read

## Inferred Memories

---

But how are these read/write operations included in your code?

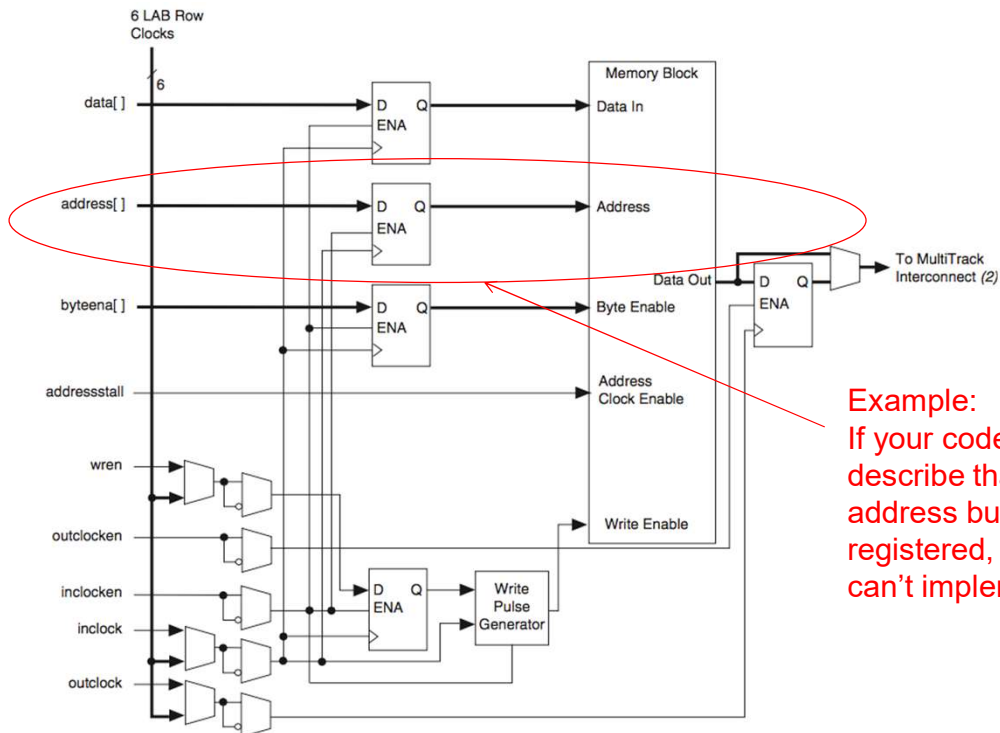
- Are they in a sequential block?
- Does there have to be any relation between read and write?
- How many read and writes are allowed?

The challenge is that the manner in which you include these operations in your code dictates the timing that Quartus Prime has to provide. The embedded RAM blocks **may not be able** to implement the timing that you describe.

- Simple example: if you specify three reads in one cycle, the embedded RAMs won't be able to do it.
- Another example: if your address is not registered, the embedded RAMS won't be able to do it.
- If it does not, Quartus will map it to flip-flop memories

Page 47

## From before: Single-Port Memory



From Altera, Cyclone II handbook

Page 48



## Inferred Memories

---

Coding rules are well documented in Quartus Prime Handbook  
(which you can find on the internet)

Intel Quartus Prime Pro Edition User Guide: Design Recommendations  
1.4.1 Inferring RAM functions from HDL Code

<https://www.intel.ca/content/www/ca/fr/programmable/documentation/sbc1513987577203.html#mwh1409959579917>

**WARNING: other tool / FPGA device vendors have different guidelines**

- Contains a number of templates that will synthesize to embedded block RAMs

We will not go through the details (read Handbook)

Page 49

In general, Quartus may find it difficult to identify embedded memories in your logic. Sometimes it will use **flip-flop memories**.

**Suggestion #1:**

Design your memory as a separate component, following the Handbook's "inferred memory" rules.

**Suggestion #2:**

If using multiple FPGA vendors (eg, Xilinx and Intel), they both have VHDL/Verilog style guides for inferring memory.

**READ THEM** and follow their advice. **When in doubt, test it out!**

**Suggestion #3:**

Use the Quartus-generated memory blocks as a last resort.

Page 50

END HERE FOR NOW

### Next topic: Scheduling

This is important because the read and write port configuration impacts how you can schedule your algorithm in hardware. In the lab, you will implement an algorithm and scheduling the operations will be a key challenge.

Page 52

## Scheduling

---

Very often, the availability of memory ports dictates scheduling.

Scheduling of an algorithm subject to memory port availability:

Two possibilities:

1. If you are using embedded memory blocks: choose single or dual-port memories. It is then *up to you* to come up with a schedule that meets the constraints on the number of ports
2. If you are using flip-flop memories, you can give it any schedule you like, and the tool will generate an appropriate number of ports
  - Very dangerous!

Page 53

## Scheduling

---

Very often, the availability of memory ports dictates scheduling.

Scheduling of an algorithm subject to memory port availability:

Two possibilities:

1. If you are using embedded memory blocks: choose single or dual-port memories. It is then *up to you* to come up with a schedule that meets the constraints on the number of ports
2. If you are using flip-flop memories, you can give it any schedule you like, and the tool will generate an appropriate number of ports  
- Very dangerous!

For the next few slides: given an algorithm, schedule it assuming we are using embedded memory with a single read or write port.

- Can do at most one read or one write each cycle

Page 55

## Scheduling

We will consider this example: (first part of ARC4 encryption):  
s is a 256x8 bit memory (256 bytes, each consisting of 8 bits):

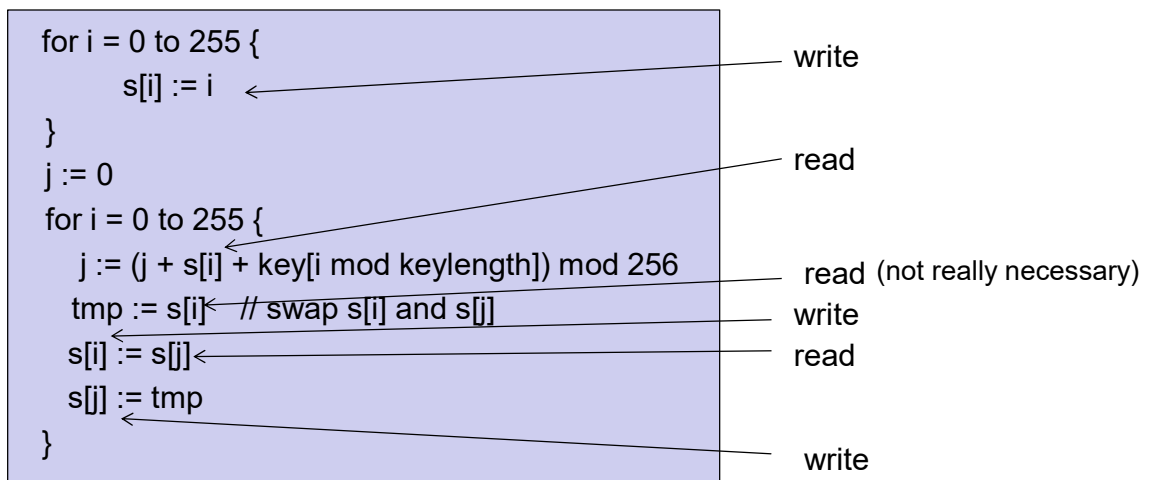
```
for i = 0 to 255 {  
    s[i] := i  
}  
j := 0  
for i = 0 to 255 {  
    j := (j + s[i] + key[i mod keylength]) mod 256  
    tmp := s[i] // swap s[i] and s[j]  
    s[i] := s[j]  
    s[j] := tmp  
}
```

Note: this is **pseudo-code**, not Verilog!

Page 56



Identify reads and writes to memory s:



```

for i = 0 to 255 {
    s[i] := i
}
j := 0
for i = 0 to 255 {
    j := (j + s[i] + key[i mod keylength]) mod 256
    tmp := s[i] // swap s[i] and s[j]
    s[i] := s[j]
    s[j] := tmp
}

```

Annotations for the code block:

- write (points to `s[i] := i`)
- read (points to `s[i]` in the second loop)
- read (unnecessary) (points to `s[i]` in the second loop)
- write (points to `s[i]` in the swap line)
- read (points to `s[j]` in the swap line)
- write (points to `s[j] := tmp`)

Start with the first for loop:

Cycle 0: write 0 to location 0  
 Cycle 1: write 1 to location 1  
 Cycle 2: write 2 to location 2  
 ...  
 Cycle 255: write 255 to location 255

Important point: we can only write to one location per cycle, so we need 256 cycles to complete this loop.

```

for i = 0 to 255 {
    s[i] := i
}
j := 0
for i = 0 to 255 {
    j := (j + s[i] + key[i mod keylength]) mod 256
    tmp := s[i] // swap s[i] and s[j]
    s[i] := s[j]
    s[j] := tmp
}

```

Labels and arrows from the right side of the code block:

- write (points to `s[i] := i`)
- read (points to `s[i]` in `j := (j + s[i] + ...)`)
- read (unnecessary) (points to `s[i]` in `tmp := s[i]`)
- write (points to `s[i] := s[j]`)
- read (points to `s[j]` in `s[i] := s[j]`)
- write (points to `s[j] := tmp`)

Now the next for loop:

Cycle 256: read s[i] (i==0)  
 Cycle 257: compute j , read s[j]  
 Cycle 258: write s[i]  
 Cycle 259: write s[j]  
 Cycle 260: read s[i] (i==1)  
 ...

This is just one schedule. There are other ways to do it, but the important point is that there is at most one memory access per cycle since there is only one port


But this schedule MAY OR MAY NOT work! Why?

To understand, let's look at typical  
memory block and timing.

It depends on whether you have  
**synchronous** or **asynchronous** data-out  
for reading memory.

(first, I'll show you the location of two problems)

Page 60



Now the next for loop...

Cycle 256: read `s[i]` ( $i=0$ )  
 Cycle 257: compute `j`, read `s[j]`  
 Cycle 258: write `s[i]`  
 Cycle 259: write `s[j]`  
 Cycle 260: read `s[i]` ( $i=1$ )  
 ...

**In Cycle 256, we read `s[i]`**  
**In Cycle 257, we hope to use**  
**the value. We can't do this if**  
**we use sync data-out!**

for i = 0 to 255 {	
s[i] := i	← write
}	
j := 0	
for i = 0 to 255 {	
j := (j + s[i] + key[i mod keylength]) mod 256	← read (unnecessary)
tmp := s[i]	← write
s[i] := s[j]	← read
s[j] := tmp	← write
}	

Now the next for loop...

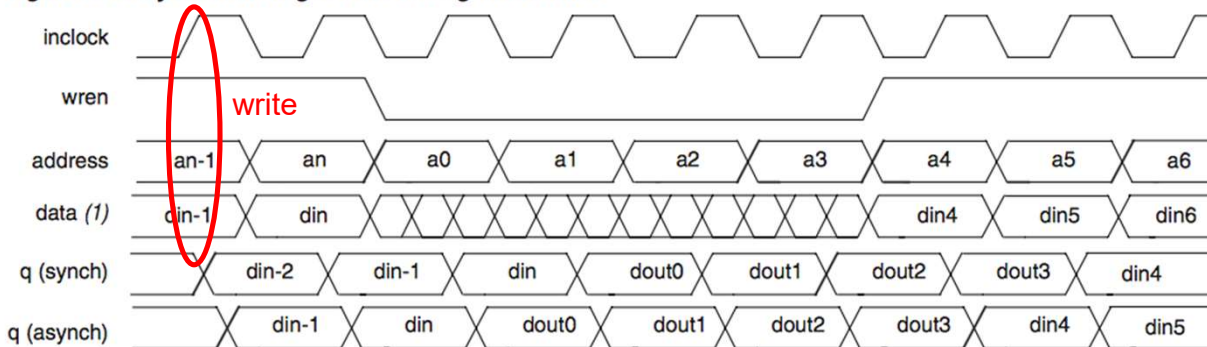
Cycle 256: read s[i] (i==0)  
 Cycle 257: compute j , read s[j]  
 Cycle 258: write s[i]  
 Cycle 259: write s[j]  
 Cycle 260: read s[i] (i==1)  
 ...

**In Cycle 257, we read s[j]**  
**In Cycle 258, we hope to use**  
**the value to write s[i]. We can't**  
**do this if we use sync data-out!**

## Memory Timing

This is a specific example: the embedded memory blocks on Cyclone II:

**Figure 8–7. Cyclone II Single-Port Timing Waveforms**

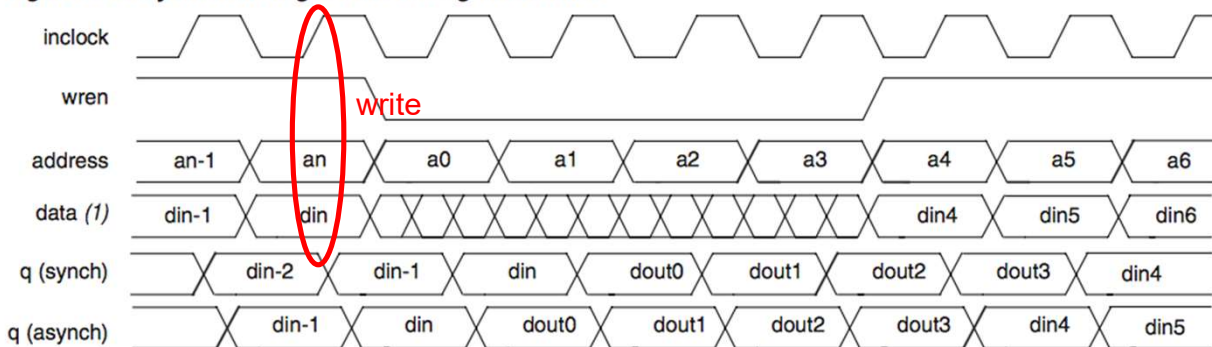


From Altera, Cyclone II handbook

## Memory Timing

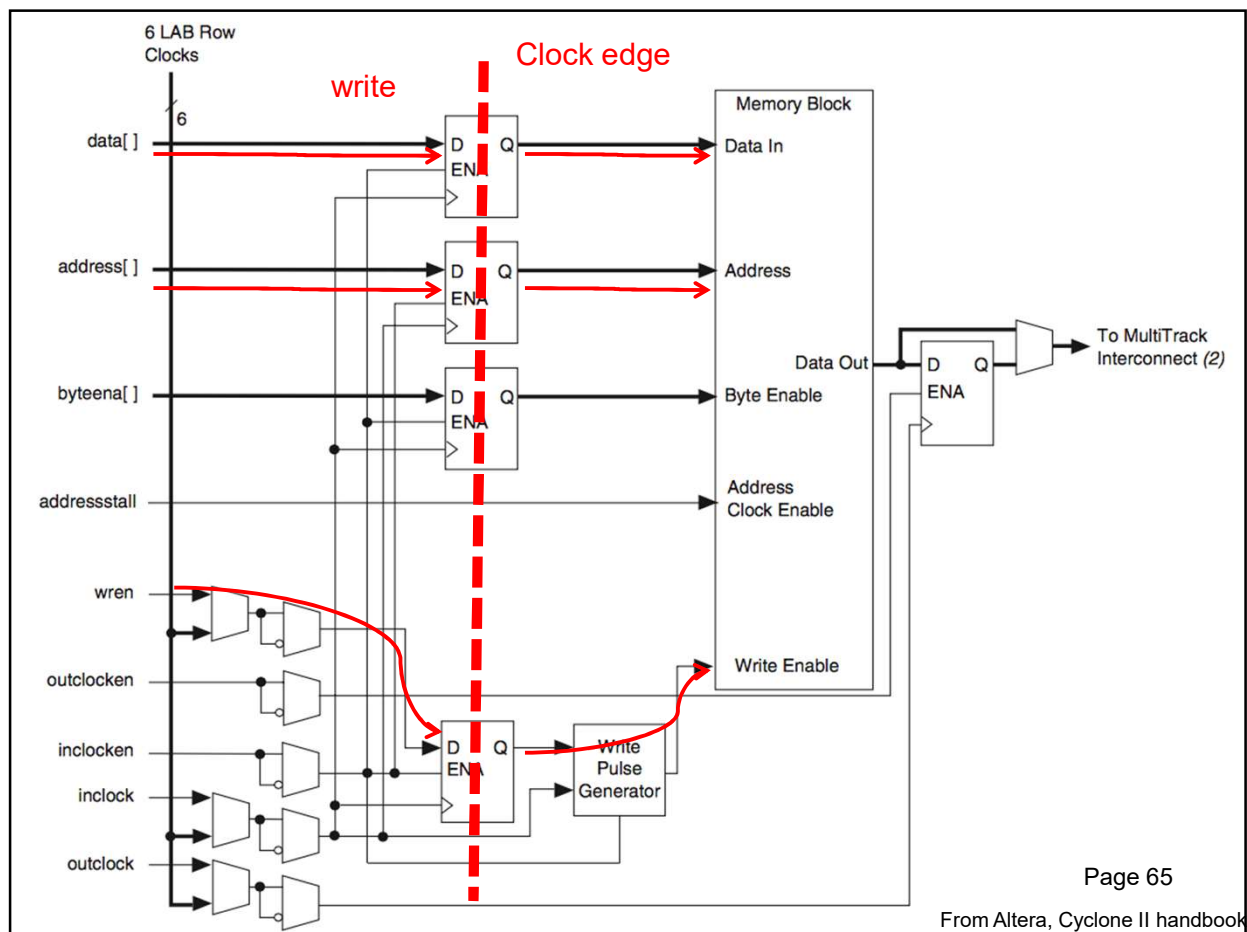
This is a specific example: the embedded memory blocks on Cyclone II:

**Figure 8–7. Cyclone II Single-Port Timing Waveforms**



From Altera, Cyclone II handbook





Page 65

From Altera, Cyclone II handbook

```

for i = 0 to 255 {
    s[i] := i
}
j := 0
for i = 0 to 255 {
    j := (j + s[i] + key[i mod keylength]) mod 256
    tmp := s[i] // swap s[i] and s[j]
    s[i] := s[j]
    s[j] := tmp
}

```

Annotations for the code block:

- write (points to `s[i] := i`)
- read (points to `s[i]` in the second loop)
- read (unnecessary) (points to `s[i]` in the second loop)
- write (points to `tmp := s[i]`)
- read (points to `s[j]` in the second loop)
- write (points to `s[j] := tmp`)

Start with the first for loop:

Cycle 0: write 0 to location 0

Cycle 1: write 1 to location 1

Cycle 2: write 2 to location 2

...

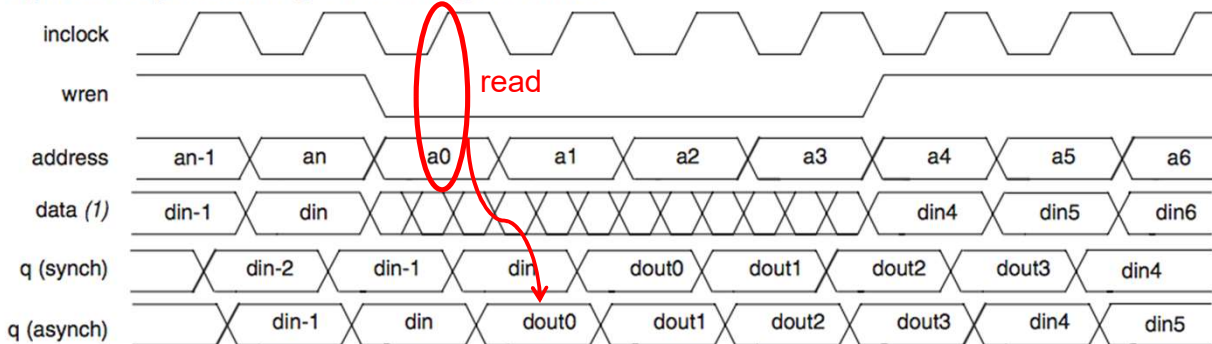
Cycle 255: write 255 to location 255

**This is all ok. We can issue writes in consecutive cycles**

## Memory Timing

This is a specific example: the embedded memory blocks on Cyclone II:

**Figure 8–7. Cyclone II Single-Port Timing Waveforms**

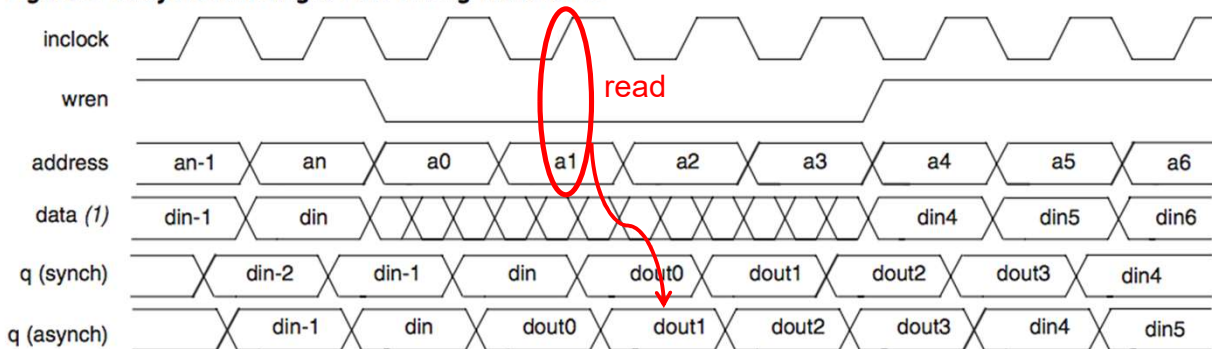


From Altera, Cyclone II handbook

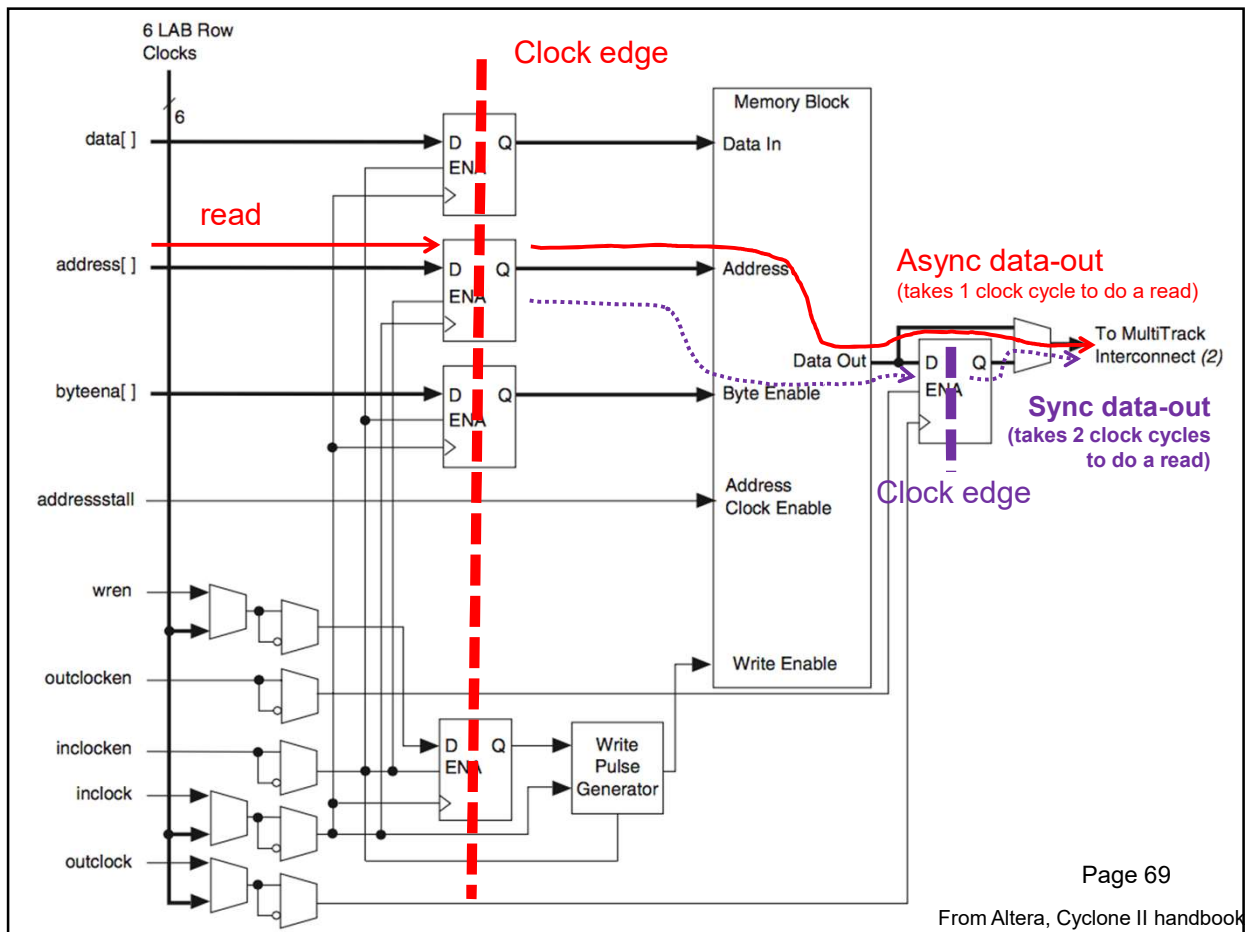
## Memory Timing

This is a specific example: the embedded memory blocks on Cyclone II:


**Figure 8–7. Cyclone II Single-Port Timing Waveforms**

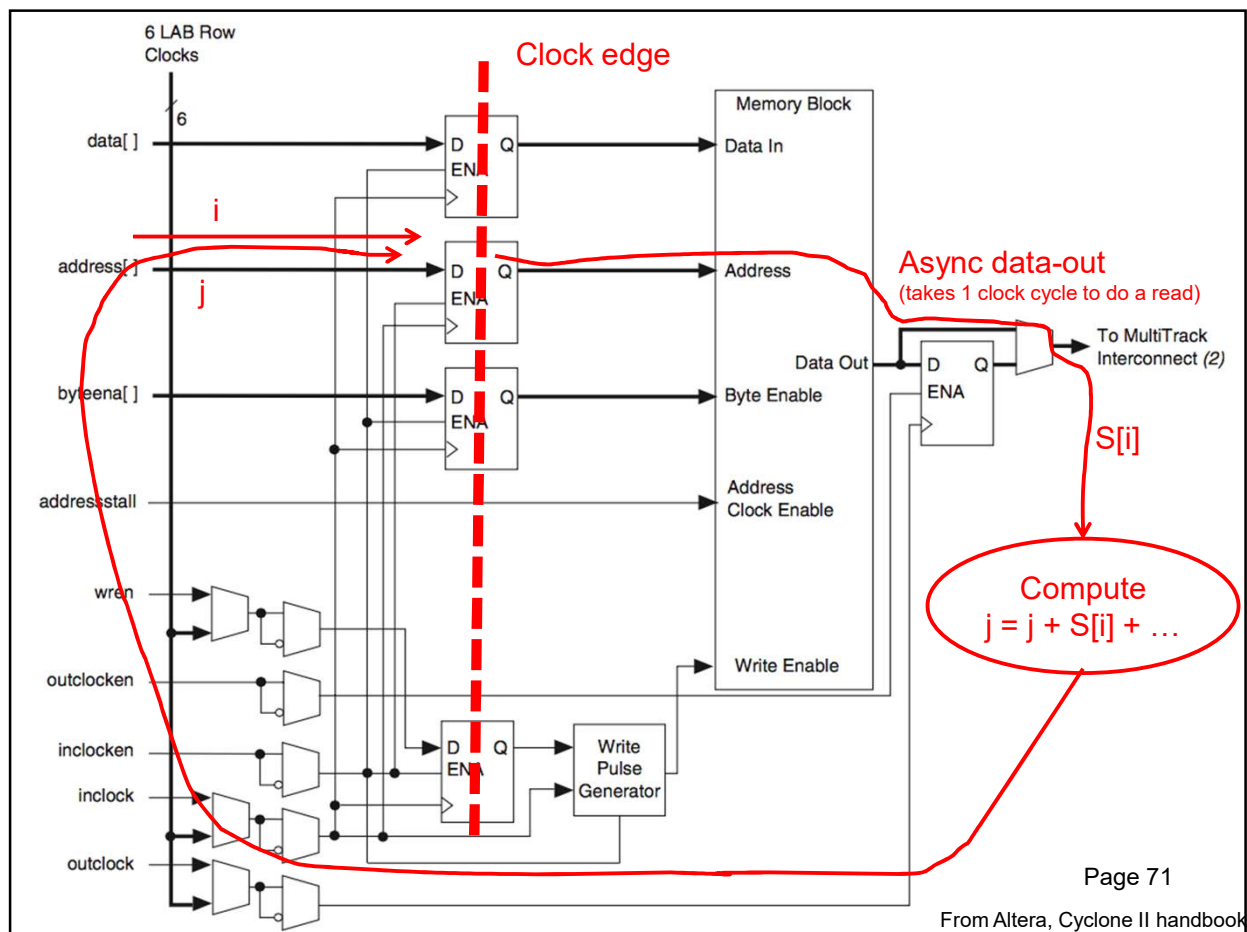


From Altera, Cyclone II handbook



## Memory Timing





Page 71

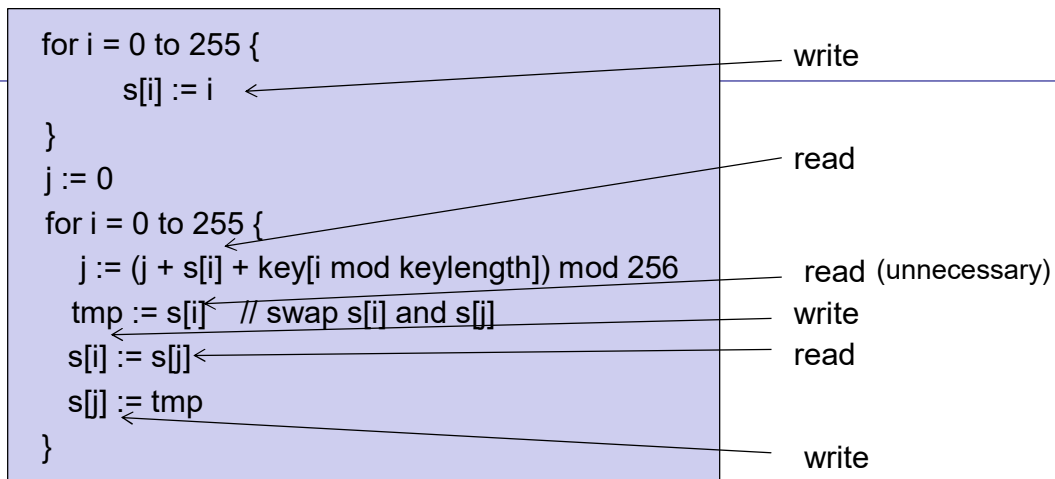
From Altera, Cyclone II handbook

So what does this mean for our example schedule?

I'll show 2 problems...

Page 72





Now the next for loop...

Cycle 256: read s[i] (i==0)  
 Cycle 257: compute j, read s[j]  
 Cycle 258: write s[i]  
 Cycle 259: write s[j]  
 Cycle 260: read s[i] (i==1)  
 ...

**In Cycle 256, we read s[i]**  
**In Cycle 257, we hope to use**  
**the value. We can't do this if**  
**we use sync data-out!**

<pre> for i = 0 to 255 {     s[i] := i } j := 0 for i = 0 to 255 {     j := (j + s[i] + key[i mod keylength]) mod 256     tmp := s[i] // swap s[i] and s[j]     s[i] := s[j]     s[j] := tmp } </pre>	<div style="margin-bottom: 10px;">← write</div> <div style="margin-bottom: 10px;">← read</div> <div style="margin-bottom: 10px;">← read (unnecessary)</div> <div style="margin-bottom: 10px;">← write</div> <div style="margin-bottom: 10px;">← read</div> <div style="margin-bottom: 10px;">← write</div>
---	--

Now the next for loop...

Cycle 256: read s[i] (i==0)  
 Cycle 257: compute j , read s[j]  
 Cycle 258: write s[i]  
 Cycle 259: write s[j]  
 Cycle 260: read s[i] (i==1)  
 ...

**In Cycle 257, we read s[j]**  
**In Cycle 258, we hope to use**  
**the value to write s[i]. We can't**  
**do this if we use sync data-out!**

## TL;DR

be very careful about modules that use memory

- can't write to same port multiple times
  - if you do and use implicit memory instances, you will get "memories" from FFs
- can't load (these) memories within one cycle
- need to design your memory client carefully

Page 75

## Learning Objectives for Slide Set

---

After this slide set, you will be able to:

1. Understand the differences, advantages and disadvantages of off-chip memory and on-chip memory
2. Understand three ways on-chip memory can be implemented in an FPGA
3. Understand the notion of memory ports, and be able to speak about single-port, dual-port and true-dual port memory configurations
4. Understand how the number and type of ports impacts overhead
5. Be able to schedule an algorithm subject to the availability of memory ports
6. Understand typical on-chip memory timing
7. Be able to write Verilog descriptions containing memory

Page 76